

Statistical analysis 12.04.17

```
# Load datasets from R
library(datasets)

# Load the library we are going to use
library(data.table)
library(psych)
library(MASS)
```

Little reminder on comparisons

Element wise comparison, compares each element of the vector.

```
a = c(1,2,3,4)

# OR is joining - use ONE |
b = a[((a>3) | (a<2))]
print(b)

## [1] 1 4

# AND is the intersection - use ONE &
b = a[((a>3) & (a<2))]
print(b)

## numeric(0)

##### Operators #####
# == equals to
# <= smaller than
# >= greater than
# != different than
# && Conditional and
# || conditional or
# / elementwise or
# & elementwise and
```

Check what your datasets

Go to : <http://127.0.0.1:8617/help/library/datasets/html/00Index.html> Or write : ?datasets then index bottom of the page

Categorical datasets

```
print(HairEyeColor)
```

```
## , , Sex = Male
##
##      Eye
## Hair   Brown Blue Hazel Green
## Black   32   11   10    3
## Brown   53   50   25   15
## Red     10   10    7    7
## Blond    3   30    5    8
##
## , , Sex = Female
##
##      Eye
## Hair   Brown Blue Hazel Green
## Black   36    9    5    2
## Brown   66   34   29   14
## Red     16    7    7    7
## Blond    4   64    5    8
```

Explore the structure of the data. Here it is a table NOT a data.table

```
print(str(HairEyeColor))
```

```
## table [1:4, 1:4, 1:2] 32 53 10 3 11 50 10 30 10 25 ...
## - attr(*, "dimnames")=List of 3
## ..$ Hair: chr [1:4] "Black" "Brown" "Red" "Blond"
## ..$ Eye : chr [1:4] "Brown" "Blue" "Hazel" "Green"
## ..$ Sex : chr [1:2] "Male" "Female"
## NULL
```

Get the marginal count for one variable: the sum of all the counts by this variable

```
eyes = margin.table(HairEyeColor, 2)
print(eyes)
```

```
## Eye
## Brown Blue Hazel Green
## 220 215 93 64
```

```
# convert to proportions
prop.table(eyes)
```

```
## Eye
##      Brown      Blue      Hazel      Green
## 0.3716216 0.3631757 0.1570946 0.1081081
```

```
# test significantly differ from H0 equal proportions
chisq.test(eyes)
```

```
##
## Chi-squared test for given probabilities
##
## data: eyes
## X-squared = 133.47, df = 3, p-value < 2.2e-16
```

```
# test with an H0 with custom population proportions (p = c(...,...))
chisq.test(eyes, p=c(.41,.32,.15,.12))
```

```
##
## Chi-squared test for given probabilities
```

```
##
## data:  eyes
## X-squared = 6.4717, df = 3, p-value = 0.09079
# Different structure for a data.table
dt = data.table(a = rep(23,4))
str(dt)

## Classes 'data.table' and 'data.frame':  4 obs. of  1 variable:
## $ a: num  23 23 23 23
## - attr(*, ".internal.selfref")=<externalptr>
```

T-Test


SAMPLE MEAN - 2 SAMPLE MEAN - 2 tail - 1 tail

```
# Using dataset mtcars
summary(mtcars)

##      mpg          cyl          disp          hp
##  Min.   :10.40   Min.   :4.000   Min.   : 71.1   Min.   : 52.0
## 1st Qu.:15.43   1st Qu.:4.000   1st Qu.:120.8   1st Qu.: 96.5
## Median :19.20   Median :6.000   Median :196.3   Median :123.0
## Mean   :20.09   Mean   :6.188   Mean   :230.7   Mean   :146.7
## 3rd Qu.:22.80   3rd Qu.:8.000   3rd Qu.:326.0   3rd Qu.:180.0
## Max.   :33.90   Max.   :8.000   Max.   :472.0   Max.   :335.0
##      drat          wt          qsec          vs
##  Min.   :2.760   Min.   :1.513   Min.   :14.50   Min.   :0.0000
## 1st Qu.:3.080   1st Qu.:2.581   1st Qu.:16.89   1st Qu.:0.0000
## Median :3.695   Median :3.325   Median :17.71   Median :0.0000
## Mean   :3.597   Mean   :3.217   Mean   :17.85   Mean   :0.4375
## 3rd Qu.:3.920   3rd Qu.:3.610   3rd Qu.:18.90   3rd Qu.:1.0000
## Max.   :4.930   Max.   :5.424   Max.   :22.90   Max.   :1.0000
##      am          gear          carb
##  Min.   :0.0000   Min.   :3.000   Min.   :1.000
## 1st Qu.:0.0000   1st Qu.:3.000   1st Qu.:2.000
## Median :0.0000   Median :4.000   Median :2.000
## Mean   :0.4062   Mean   :3.688   Mean   :2.812
## 3rd Qu.:1.0000   3rd Qu.:4.000   3rd Qu.:4.000
## Max.   :1.0000   Max.   :5.000   Max.   :8.000

# extract mpg
hist(mtcars[,1])
```

Histogram of mtcars[, 1]




```
# test significantly different than 20 -- get confidence interval of 0.8  
t.test(mtcars[,1], mu = 20, conf.level = 0.8)
```

```
##  
## One Sample t-test  
##  
## data: mtcars[, 1]  
## t = 0.08506, df = 31, p-value = 0.9328  
## alternative hypothesis: true mean is not equal to 20  
## 80 percent confidence interval:  
## 18.69549 21.48576  
## sample estimates:  
## mean of x  
## 20.09062
```

```
# create fake data  
# Sample 18000 scores from a normal distribution of mean 5 and sd 2  
memory_test = rnorm(18000, mean = 5, sd = 2)  
  
# take only the values between 0 and 10 (you could use also the function ifelse())  
memory_test_filtered = memory_test[memory_test>0 & memory_test<10]  
hist(memory_test_filtered)
```

Histogram of memory_test_filtered




```
# test significantly different than 5  
t.test(memory_test_filtered, mu=5)
```

```
##  
## One Sample t-test  
##  
## data: memory_test_filtered  
## t = -0.74556, df = 17782, p-value = 0.4559  
## alternative hypothesis: true mean is not equal to 5  
## 95 percent confidence interval:  
## 4.961423 5.017317  
## sample estimates:  
## mean of x  
## 4.98937
```

```
# sample the same amount of subjects this time centered on 6  
memory_test_female = rnorm(18000, mean = 6, sd = 2)  
hist(memory_test_female)
```

Histogram of memory_test_female




```
describe(memory_test_female)
```

```
##      vars      n mean  sd median trimmed  mad   min   max range skew
## X1      1 18000 6.01 2.01   6.02    6.02 2.01 -1.86 15.87 17.74    0
##      kurtosis   se
## X1      -0.01 0.01
```

```
# Filter it in the same way take only values between 0 and 10
```

```
memory_test_female_filtered = memory_test_female[memory_test_female>0 & memory_test_female<10]
hist(memory_test_female_filtered)
```

Histogram of memory_test_female_filtered



```
# One tailed (alternative / alt = "greater")
t.test(memory_test_female_filtered, memory_test_filtered, alt = "greater")
```

```
##
## Welch Two Sample t-test
##
## data: memory_test_female_filtered and memory_test_filtered
## t = 45.685, df = 35349, p-value < 2.2e-16
## alternative hypothesis: true difference in means is greater than 0
## 95 percent confidence interval:
##  0.8859634      Inf
## sample estimates:
## mean of x mean of y
##  5.908424  4.989370
```

PAIRED T-TEST

T1 -> Intervention -> T2 (repeated measure)

```
## GROUP T1 - first test before the intervention
# sample fake data centered on 2
memory_test = rnorm(18000, mean = 2, sd = 1)


# a lot of the data is below 0
# but since we are going to use a paired t-test we want to keep the same amount of
# subjects and not delete them

# We need to REPLACE the values for subjects below 0
# If we replace all the values below 0 by 0 it will strongly alter the shape of the distribution
```

```
memory_test[memory_test<0] = 0
```

```
# look at the first bin on the left is a lot higher than its symmetric bin on the other side  
# The distribution is skewed, which is not bad in itself, you just have to account for it by using  
# robust statistics if the skeweness is too important  
hist(memory_test)
```

Histogram of memory_test



```
# Another way, to smooth your distribution a bit could be to replace by 0 + a random value between 0 and 1  
# (this is fake data -- not something you will want to do with real data)
```

```
memory_test = rnorm(18000, mean = 2, sd = 1)
```


```
# get the number of subjects below 0  
size_below_zero = length(memory_test[memory_test<0])
```

```
# get a vector of values between 0 and 3  
sampled_vector = seq(from = 0, to = 3, by = 0.01)
```

```
# sample from that vector the same amount of time that there are subject below 0 in the memory test  
sample_for_replacement = sample(sampled_vector, size = size_below_zero, replace = T)
```


```
# replace the values  
memory_test[memory_test<0] = sample_for_replacement  
hist(memory_test)
```


Histogram of memory_test



```
## GROUP T2 cw in between !  
# We create the second vector by adding a random number sampled from a normal distribution  
# with mean of 3 and sd of 1  
added_vector = rnorm(length(memory_test), mean = 3, sd = 1)  
hist(added_vector)
```

Histogram of added_vector




```

# our intervention has a positive effect so we add this positive vector
memory_test_t2 = copy(memory_test) + added_vector

# we build our data.table with the two samples
pairs = data.table(t1 = memory_test, t2 = memory_test_t2)

# Check visually evolution of the 100 first subjects (parcoord is part of the MASS package)
parcoord(pairs[0:100,], var.label = T)

```



```

# Paired T test with H0 difference is not significantly different than 0
t.test(memory_test_t2, memory_test, paired = T)

```

```

##
## Paired t-test
##
## data: memory_test_t2 and memory_test
## t = 402.19, df = 17999, p-value < 2.2e-16
## alternative hypothesis: true difference in means is not equal to 0
## 95 percent confidence interval:
##  2.978015 3.007185
## sample estimates:
## mean of the differences
##                2.9926

```

ANOVA


One variable

```

# Check that the number of days of holidays between countries is different
chinese = rnorm(n = 150, mean = 3, sd = 1)
japan = rnorm(n = 30, mean = 6, sd = 1)
danemark = rnorm(n = 20, mean = 30, sd = 4)
usa = rnorm(n = 400, mean = 4, sd = 2)

```

```
boxplot(chinese, japan, danemark, usa)
```



```
# We need an indicator variable, we build it by hand
group = c(rep('chinese', 150), rep('japanese', 30), rep('danish', 20), rep('american', 400))
```

```
# This is our value variable
values = c(chinese, japan, danemark, usa)
```

```
# We create a data.table with those two corresponding vector
dt = data.table(g = group, v = values)
print(dt)
```

```
##           g           v
##  1:  chinese 2.737202
##  2:  chinese 3.197161
##  3:  chinese 3.972587
##  4:  chinese 3.666251
##  5:  chinese 3.040477
## ---
## 596: american 3.228802
## 597: american 3.966468
## 598: american 3.264032
## 599: american 6.133591
## 600: american 8.025624
```

```
# Now we can perform our anova test
aov_results = aov(dt, formula = v ~ g)
summary(aov_results)
```

```
##           Df Sum Sq Mean Sq F value Pr(>F)
## g           3   15498     5166   1342 <2e-16 ***
## Residuals   596    2294         4
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```

# We can also test group difference one to one using Tukey
# Here they are all significant
tukey_results = TukeyHSD(aov_results)
print(tukey_results)

##      Tukey multiple comparisons of means
##      95% family-wise confidence level
##
## Fit: aov(formula = v ~ g, data = dt)
##
## $g
##              diff              lwr              upr      p adj
## chinese-american -0.9795362 -1.4634467 -0.4956257 1.5e-06
## danish-american  27.9429117  26.7848305  29.1009929 0.0e+00
## japanese-american  1.8346876   0.8779277   2.7914475 6.0e-06
## danish-chinese    28.9224479  27.7192885  30.1256073 0.0e+00
## japanese-chinese   2.8142238   1.8033677   3.8250799 0.0e+00
## japanese-danish  -26.1082241 -27.5672692 -24.6491790 0.0e+00

# To use the function stack() the number of observation need to be similar
chinese = rnorm(n = 50, mean = 3, sd = 1)
japan = rnorm(n = 50, mean = 6, sd = 1)
danemark = rnorm(n = 50, mean = 30, sd = 4)
usa = rnorm(n = 50, mean = 4, sd = 2)
stack_variable = stack(data.frame(cbind(chinese, japan, danemark, usa)))
result = aov(stack_variable, formula = values ~ ind)
summary(result)

##              Df Sum Sq Mean Sq F value Pr(>F)
## ind              3  24658    8219   1524 <2e-16 ***
## Residuals       196   1057         5
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

```